

# IPC: How speed is my Virtual Machine

Janardhana Reddy Naredula, Flipkart

## Abstract

Virtual machines in the cloud typically run existing general-purpose traditional operating systems(OS) such as Linux. The speed at which the instructions are executed by the CPU depends on various factors like memory access speed,locks,syscalls,hyperthreading..etc. This paper describes the details of all the parameters governing the speed of cpu. Linux and Jiny kernel[1] are used to describe these parameters. Speed of cpu can be measured in the units of Instruction Per Cycle(IPC), higher value of IPC means better speed.

## 1 Introduction

IPC depends on various parameters. IPC varies from approx 0.01 to 4.50. The following are different type of workloads, These workload are very homogenous, In the actual application workload will be heterogeneous containing combination of all the below workload:

1. **Memory Intensive:** writing or reading large amount of memory causes the IPC to slowdown. Here it depends on the size of the memory and page size. If the page size is bigger and memory is smaller, the speed will be better. Speed is proportional to (pagesize/memory size accessed).
2. **Syscall Intensive:** During syscall, some of costly instruction are executed, due to this IPC will slows down.
3. **Spinlock:** locking of the memory bus causes the IPC to slow down.
4. **Hyperthread:** Hyper Threading efficiency depends on type of workload on the logical cpus of a physical core. If all logical threads executes cpu intensive workloads then the efficiency will falls down. If the logical threads executes memory intensive or similar loads, then the instructions from logical threads will get interleaved and get the maximum efficiency.
5. **Interrupts:** TODO
6. **Network packet process:** TODO
7. **Multicore, IPI,Hlt :** TODO

**Applications on virtual machine:** The workloads described above and corresponding tests in the next section are very homogeneous in nature, means the test application stresses on a particular workload. But the practical application contain mix of all the above workloads of

various degree, due to this IPC of an realtime applications will be combinations of all the workloads mentioned above.

## 2. Test Setup:

Perf tool[2] is used to to measure IPC on the host. Vm are created of different os's like linux and Jiny kernel[1]. Jiny kernel is used in all the below tests, since it as features to disable/enable lot of OS features for the performance tests. The results will be similar to linux vm. Stress test app like "stress" is used to generate various loads like memory,syscall,.. etc inside the vm. and at the same time perf will be used to measure the IPC for the specified vm on physical host. IPC value mentioned in the below test is for entire VM not the test application.

1. Perf: To measure the speed(i.e ipc) of vm(Jiny kernel/linux) in terms of Instruction Per Cycle (IPC).
2. Application inside the Jiny kernel to generate various work loads:
  - a. Memory intensive test using stress[3]: memory size from 128k to 500M is used.
  - b. Syscall intensive test using stress[3]: calling syscall in a tightloop, and the syscall does very minimal work inside the the os.
  - c. Lock intensive test using Semaphore[5]: Two threads in tight loop acquire the lock and unlock.

### 2.1 Tests:

The speed depends on OS overhead, hardware speed, hypervisor overhead, and the actual app. For the same app, the speed between different os's/hardware will be relatively similar, means speed is mostly driven by the application logic and os configuration.

IPC is calculated from host using linux perf tool as below:

**perf kvm --guest stat -e cycles,instructions -p <pid of vm>**

Tes t#		Description ( commands = speed)	Speed (IPC)	Reason
1		Memory/cpu test with 250 bytes without syscalls:  ./stress --vm 1 --vm-bytes 250 --vm-keep=4.60	4.60	Baseline for all test, it is best speed in all the test. It does not have any memory,syscall or cpu overheads.
2	me mor y	memory 4k pages : touching the memory in tight loop. ./stress --vm 1 --vm-bytes 128k --vm-keep=0.98 ./stress --vm 1 --vm-bytes 30M --vm-keep=0.16 ./stress --vm 1 --vm-bytes 128M --vm-keep=0.15 ./stress --vm 1 --vm-bytes 500M --vm-keep = 0.14	0.12-0.16	As the memory size increases, IPC decreases.

		<code>./stress --vm 1 --vm-bytes 1500M --vm-keep=0.12</code>		
3		Memory Hugepages: <code>./stress --vm 1 --vm-bytes 30M --vm-keep=0.19</code> <code>./stress --vm 1 --vm-bytes 128M --vm-keep=0.18</code> <code>./stress --vm 1 --vm-bytes 500M --vm-keep = 0.18</code> <code>./stress --vm 1 --vm-bytes 1500M --vm-keep=0.18</code>	0.18	Gain because of less ptes, due to less TLB misses.
4		Hugepages (1 process) + cpu intensive (2nd process) <code>./stress --vm 1 --vm-bytes 1500M --vm-keep</code> + <code>./stress --vm 1 --vm-bytes 256 --vm-keep</code>	2.40	This is better than previous test, because 2nd process run at high speed, so the average of two improves.
5	sys call s	syscallintensive : 1 process	0.59	This is less when compare to the baseline because of the context switch at every syscall. Here the syscall is almost empty, it does not consume much inside the os.
6		syscallintensive : 2 or more identical process	0.49	It is less when compare to above because of extra memory pressure during context switch.
7	spi nlo ck	Semaphore and spinlock test: system calls are very minimum, two threads are in tight spin to acquire the lock. <code>./sem4 100000000 100000000000 =1.16</code>	1.16	It is less when compare to the baseline because it does the tight spin with neighbouring thread.
8	Hy pert hre adi ng imp act	Two vcpus from different vm's on the different core = 4.6 (same as test1) <code>./stress --vm 1 --vm-bytes 250 --vm-keep=4.60</code>  Two vcpus from different vm's on the same core, both are cpu intensive: <code>./stress --vm 1 --vm-bytes 250 --vm-keep=2.3</code>	2.6 to 4.6	50% loss because guest os is unaware of hyperthreading.

### 3.0 Summary:

Area	Difference in speed	Reason
------	---------------------	--------

4K versus Hugepages	$0.12 \text{ to } 0.18 = 0.06$	TLB lookup is fast in hugepages, due to this speed is higher.
Syscall overhead	$4.60 \text{ to } 0.59 = 4.10$	Switching from userspace to kernel. Slow down because of the following reasons: 1) context switch from userspace to kernel 2) when there are more than 1 processes, then additional slow down due to memory switch .
Spinlock overhead	$4.60 \text{ to } 0.83 = 3.80$	Special instruction in implementing spin lock that lock the memory bus, due to this there is slow down.
Hyperthread friendliness	2.3(different cores) to 4.6 (same core) = 2.3	vcpu will be slowed down by 50% if the vcpu of different vm's falls inside the same core, assuming both the apps are cpu intensive loads. Means the app can run twice the speed if vcpu is alone inside the physical core and the load is cpu intensive.
Interrupts, IPI, Multicore, Networking	TODO	TODO

## References:

- [1] Jiny Kernel: <https://github.com/naredula-jana/Jiny-Kernel>
- [2] Perf: source code inside linux kernel repo
- [3] stress test tool: <http://people.seas.harvard.edu/~apw/stress/>
- [4] Paper from google: <http://www.e-wilkes.com/john/papers/2013-EuroSys-CPI2.pdf>
- [5] semaphore test: